

# Serverless Architectures with Amazon DynamoDB and Amazon Kinesis Streams with AWS Lambda

Krishnasis Mandal

Associate Analyst

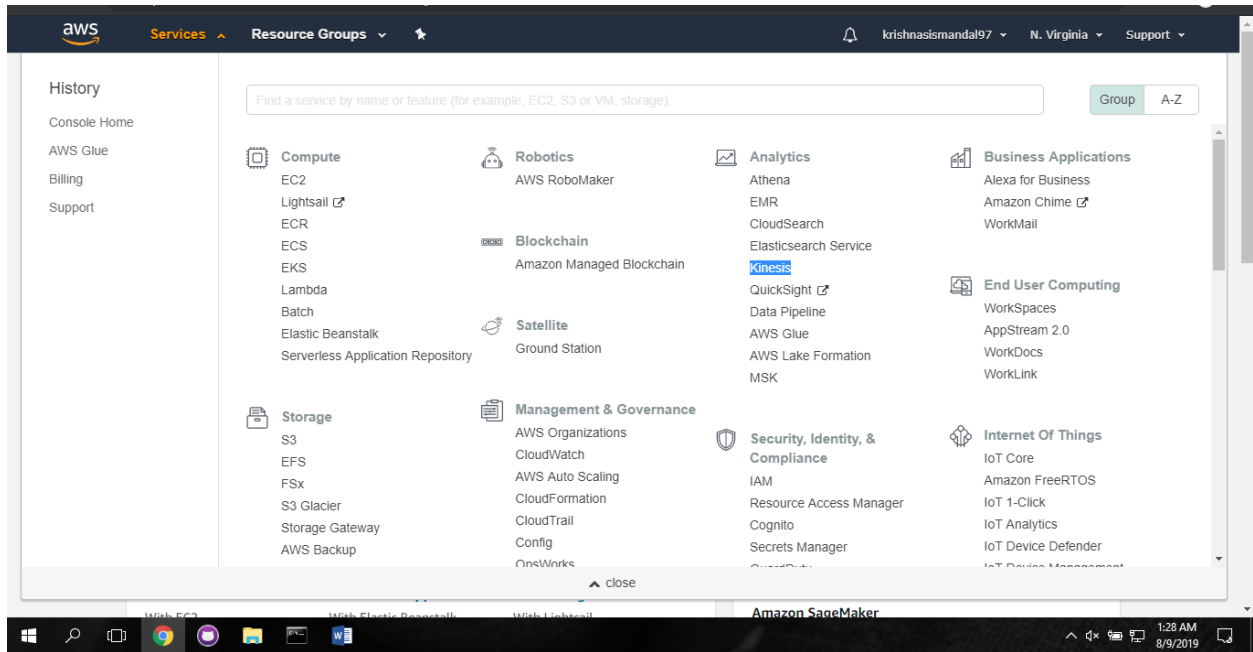


## Contents

Creating an Amazon Kinesis Stream .....	3
Creating a Lambda Function .....	5
Testing the Lambda Function.....	8
Creating Tables in DynamoDB.....	10
Create a Lambda Function .....	12

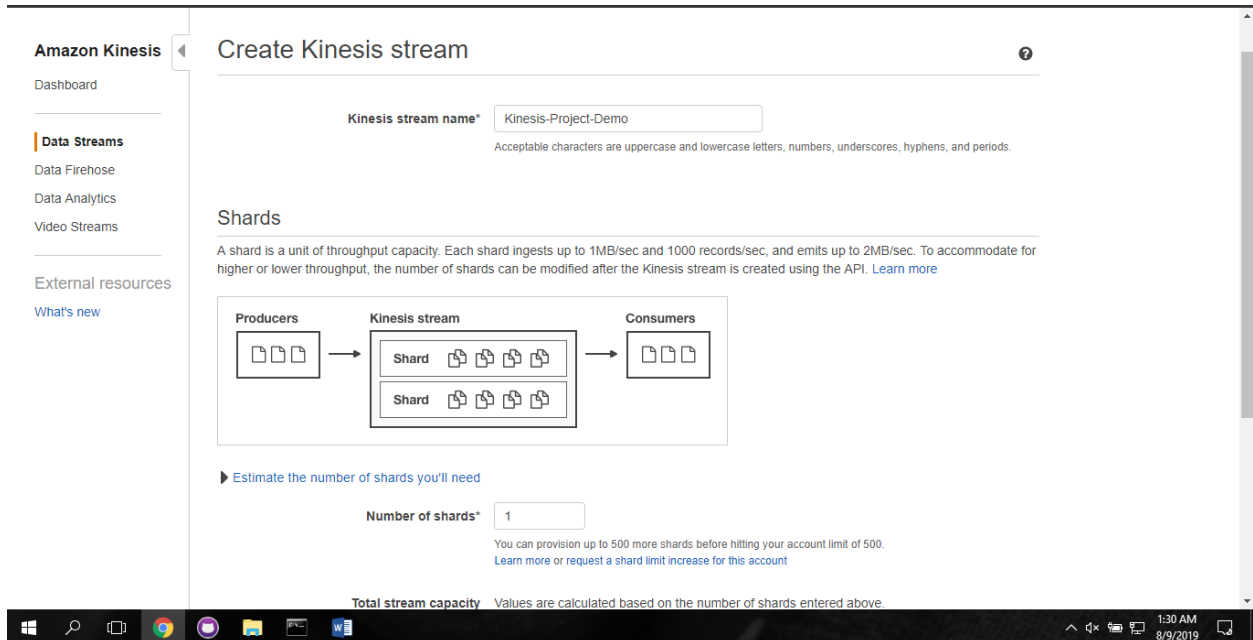
# Creating an Amazon Kinesis Stream

Navigate to Kinesis Service from the AWS Management Console



Click on Get Started

Click on Create Data Stream



Cancel

Create Kinesis stream

Click on Create Kinesis Stream

The Kinesis Stream will be created.

The screenshot displays the Amazon Kinesis console interface. On the left, a sidebar menu includes 'Amazon Kinesis', 'Dashboard', 'Data Streams' (highlighted), 'Data Firehose', 'Data Analytics', 'Video Streams', 'External resources', and 'What's new'. The main content area is titled 'Kinesis streams' and contains a description of Kinesis data streams. It shows 'Total shards in use: 1' and 'Total shards remaining: 499'. A green success message states 'Stream Kinesis-Project-Demo has been created'. Below this, there are buttons for 'Create Kinesis stream', 'Connect Kinesis consumers', and 'Actions'. A search bar and pagination controls are also present. A table lists the created stream:

<input type="checkbox"/>	Kinesis stream name	Number of shards	Status	Consumers using enhanced fan-out
<input type="checkbox"/>	Kinesis-Project-Demo	1	Active	0

The footer of the console shows 'Feedback', 'English (US)', copyright information, and links to 'Privacy Policy' and 'Terms of Use'. The Windows taskbar at the bottom indicates the time is 1:31 AM on 8/9/2019.

# Creating a Lambda Function

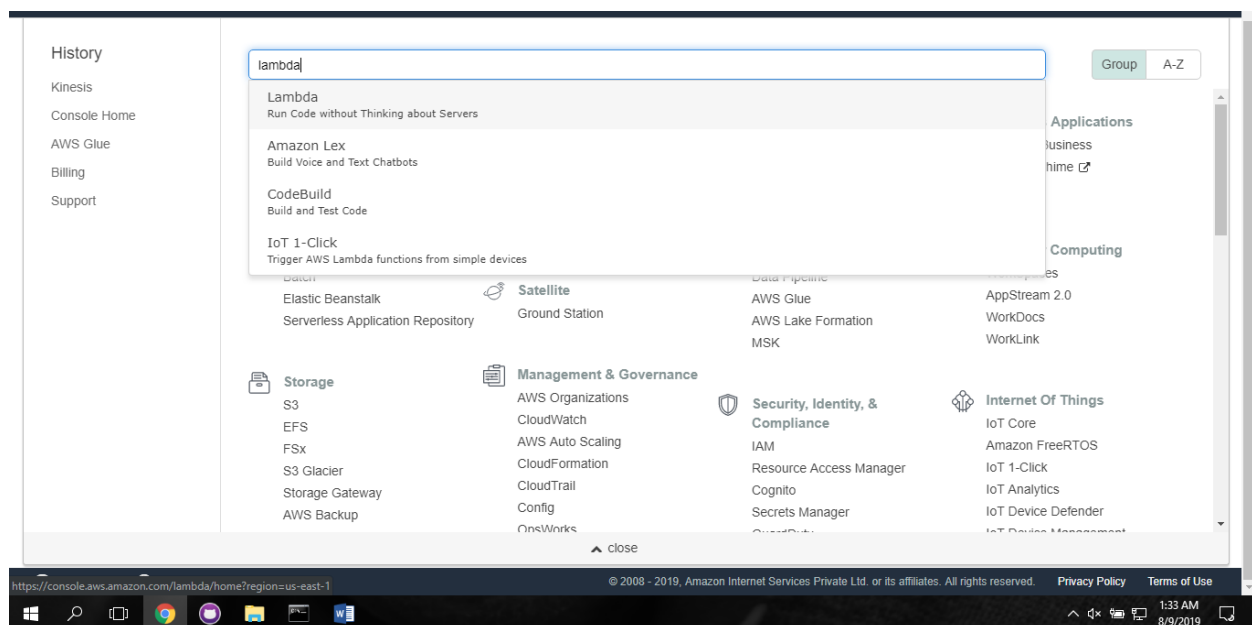
Create a IAM Role for Lambda with the Policies – AmazonKinesisFullAccess, CloudWatchFullAccess, AWSLambdaFullAccess, selected.

The screenshot shows the 'Create role' page in the AWS IAM console, specifically the 'Review' step. The page has a progress indicator at the top right with four steps, where the fourth step is highlighted. The 'Review' section contains the following information:

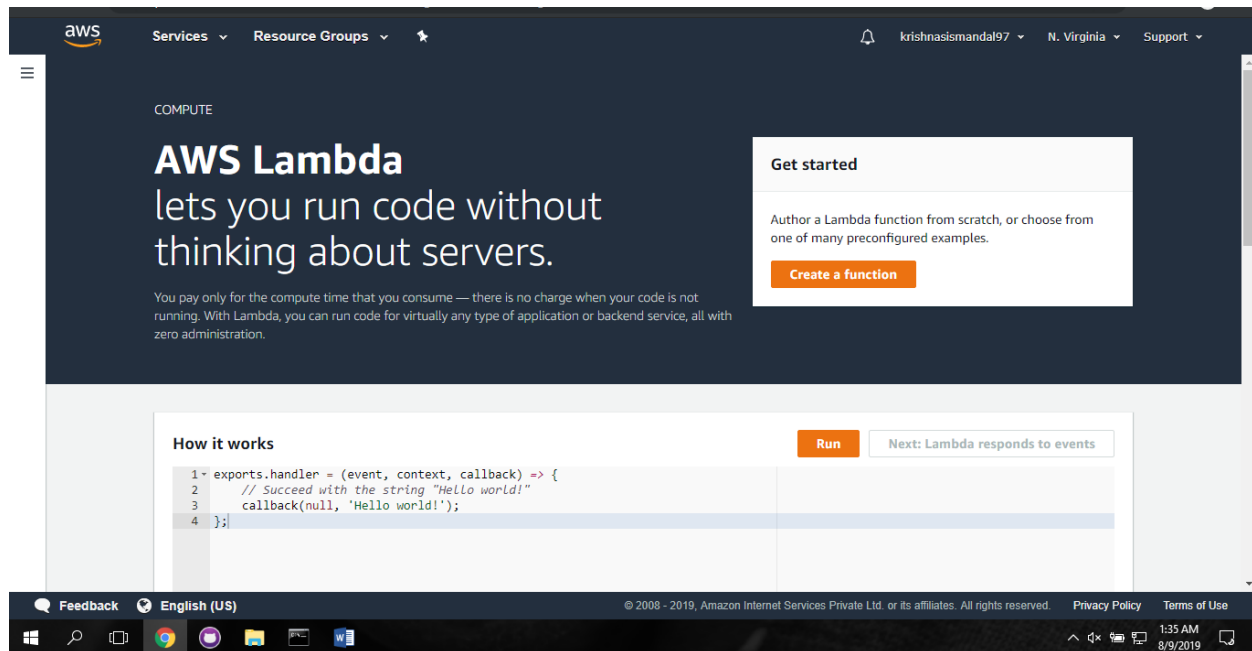
- Role name:** LambdaRole. Below the input field, it says: 'Use alphanumeric and '+=, @, \_' characters. Maximum 64 characters.'
- Role description:** Allows Lambda functions to call AWS services on your behalf. Below the input field, it says: 'Maximum 1000 characters. Use alphanumeric and '+=, @, \_' characters.'
- Trusted entities:** AWS service: lambda.amazonaws.com
- Policies:** Three policies are listed with checkboxes: AmazonKinesisFullAccess, CloudWatchFullAccess, and AWSLambdaFullAccess. All three are checked.
- Permissions boundary:** Permissions boundary is not set.

At the bottom of the form, there are buttons for 'Cancel', 'Previous', and 'Create role'. Below the form, there is a footer with 'Feedback', 'English (US)', copyright information, 'Privacy Policy', and 'Terms of Use'. The Windows taskbar is visible at the very bottom.

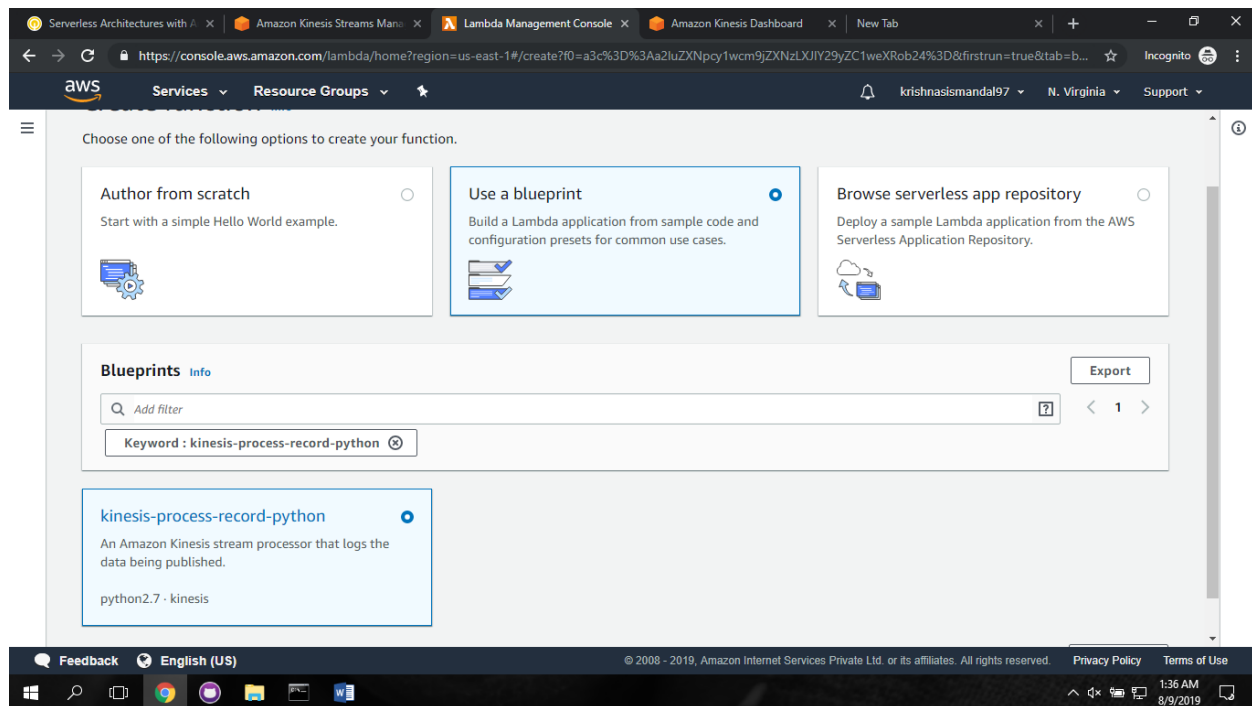
Navigate to AWS Lambda from the Services menu of the AWS Management Console.



Click on Create a Function



Click on Use a blueprint, and then search for kinesis-process-record-python.



Scroll down and click on Configure.

**Lambda > Functions > Create function > Using blueprint kinesis-process-record-python**

**Basic information** [Info](#)

Function name

Execution role  
Choose a role that defines the permissions of your function. To create a custom role, go to the [IAM console](#).

Existing role  
Choose an existing role that you've created to be used with this Lambda function. The role must have permission to upload logs to Amazon CloudWatch Logs.  
 [View the LambdaRole role on the IAM console.](#)

**Kinesis trigger** [Remove](#)

Kinesis stream  
Select a Kinesis stream to listen for updates on.

Feedback English (US) © 2008 - 2019, Amazon Internet Services Private Ltd. or its affiliates. All rights reserved. Privacy Policy Terms of Use 1:40 AM 8/9/2019

**Kinesis stream**  
Select a Kinesis stream to listen for updates on.

**Consumer**  
Select an optional [consumer](#) of your stream to listen for updates on.

**Batch size**  
The largest number of records that will be read from your stream at once.

**Starting position**  
The position in the stream to start reading from. For more information, see [ShardIteratorType](#) in the Amazon Kinesis API Reference.

In order to read from the Kinesis trigger, your execution role must have proper permissions.

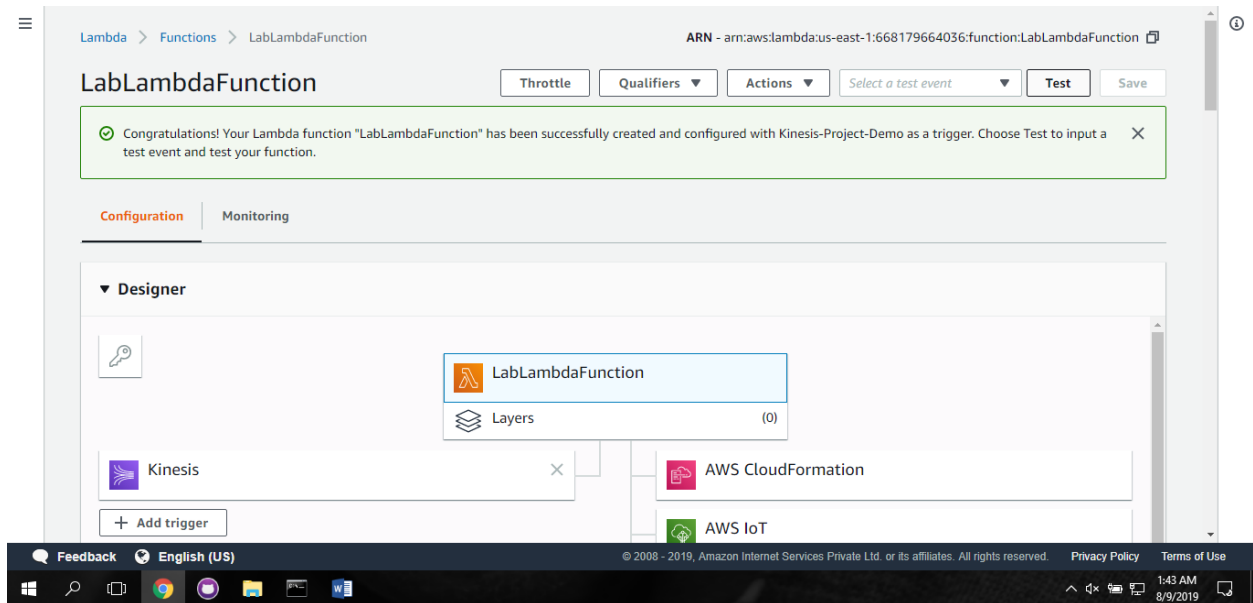
☒ **Enable trigger**  
Enable the trigger now, or create it in a disabled state for testing (recommended).

**Lambda function code**

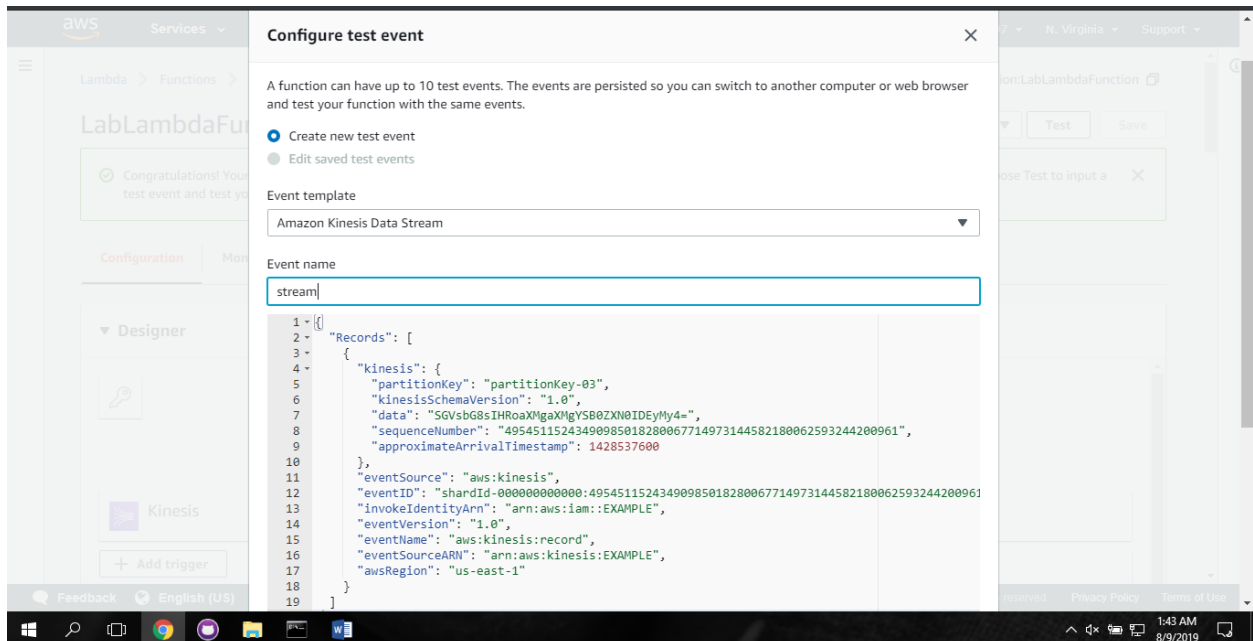
Feedback English (US) © 2008 - 2019, Amazon Internet Services Private Ltd. or its affiliates. All rights reserved. Privacy Policy Terms of Use 1:40 AM 8/9/2019

Select the appropriate parameters, review the code, and click on Create Function.

# Testing the Lambda Function

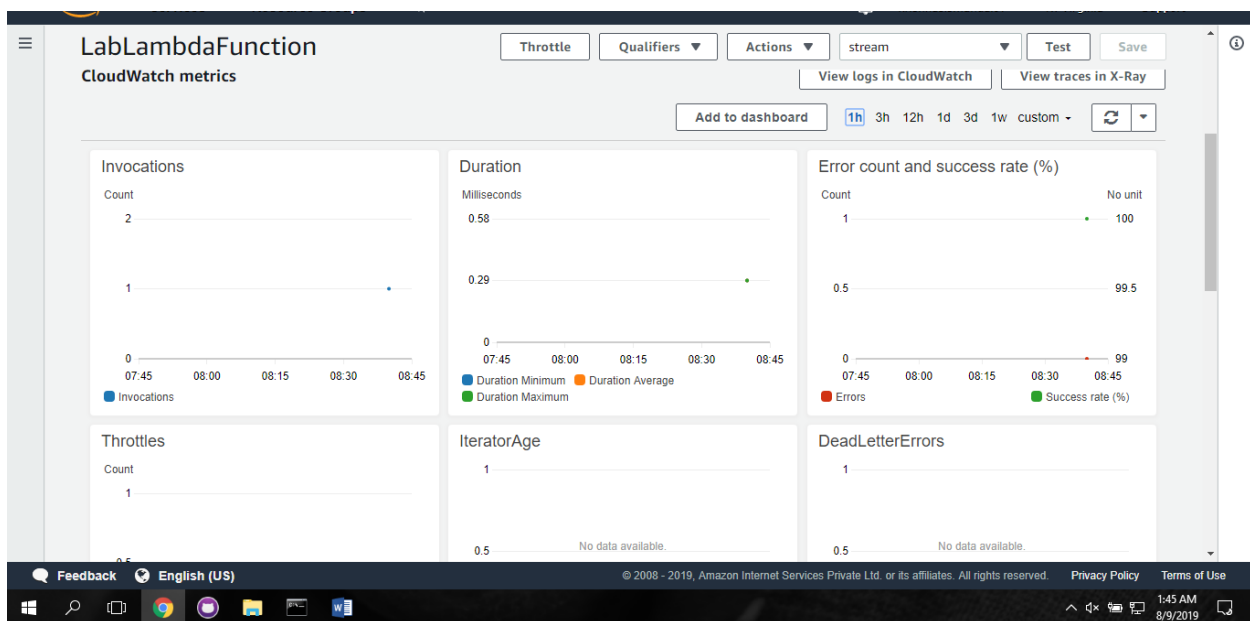
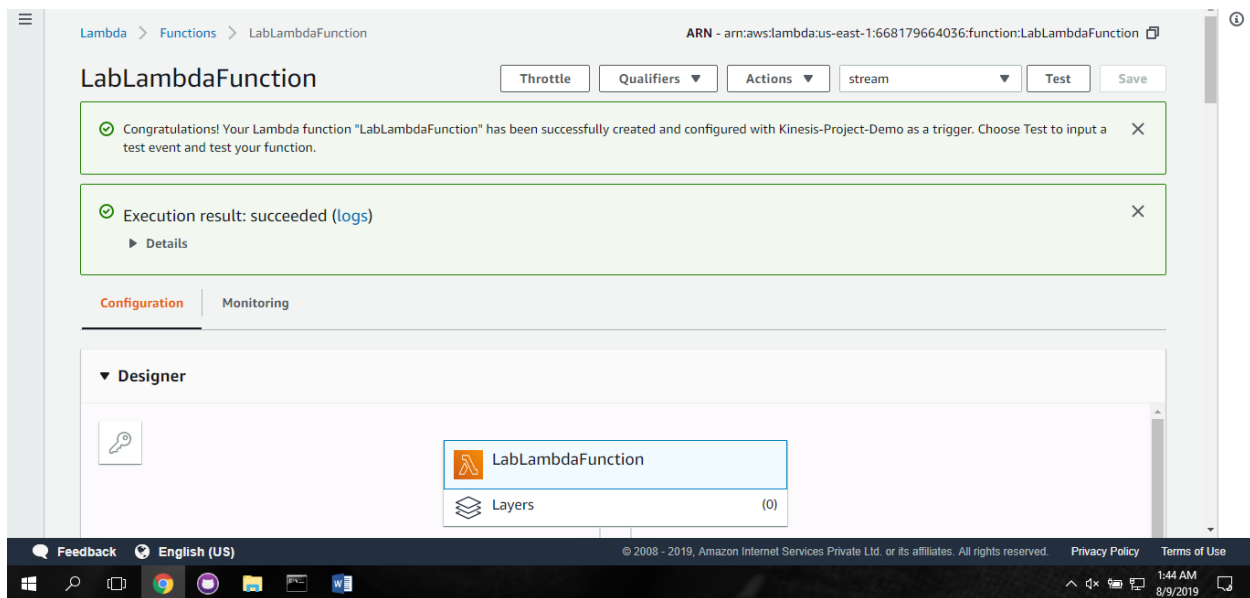


Click on Test button.



Click on Create and the Test button again.





Click on Monitoring to get the detailed CloudWatch reports pertaining to the Test function.

# Creating Tables in DynamoDB

The screenshot shows the 'Create DynamoDB table' page in the AWS Management Console. The page title is 'Create DynamoDB table' with a 'Tutorial' link. Below the title, a brief description of DynamoDB is provided. The 'Table name\*' field is set to 'GameScoreRecords'. The 'Primary key\*' section shows 'Partition key' as 'RecordID' with a 'Number' data type. There is an 'Add sort key' checkbox which is unchecked. Under 'Table settings', the 'Use default settings' checkbox is checked. A list of default settings is shown: 'No secondary indexes', 'Provisioned capacity set to 5 reads and 5 writes', 'Basic alarms with 80% upper threshold using SNS topic "dynamodb"', and 'Encryption at Rest with DEFAULT encryption type'. A blue information box states: 'You do not have the required role to enable Auto Scaling by default. Please refer to documentation.' The footer includes 'Feedback', 'English (US)', copyright information, 'Privacy Policy', 'Terms of Use', and a system clock showing 1:50 AM on 8/9/2019.

Create DynamoDB table [Tutorial](#)

DynamoDB is a schema-less database that only requires a table name and primary key. The table's primary key is made up of one or two attributes that uniquely identify items, partition the data, and sort data within each partition.

Table name\*

Primary key\* Partition key

☐ Add sort key

Table settings

Default settings provide the fastest way to get started with your table. You can modify these default settings now or after your table has been created.

☒ Use default settings

- No secondary indexes.
- Provisioned capacity set to 5 reads and 5 writes.
- Basic alarms with 80% upper threshold using SNS topic "dynamodb".
- Encryption at Rest with DEFAULT encryption type.

**You do not have the required role to enable Auto Scaling by default.**  
Please refer to [documentation](#).

Feedback English (US) © 2008 - 2019, Amazon Internet Services Private Ltd. or its affiliates. All rights reserved. Privacy Policy Terms of Use 1:50 AM 8/9/2019

Create a DynamoDB table with the appropriate parameters.

Create another DynamoDB table.

The screenshot shows the 'GameScoresByUser' table details page in the AWS Management Console. The page title is 'GameScoresByUser' with a 'Close' link. The 'Overview' tab is selected. A blue banner at the top says 'Table is being created'. Below this, the 'Recent alerts' section shows 'No CloudWatch alarms have been triggered for this table.' The 'Stream details' section shows 'Stream enabled' as 'No', 'View type' as '-', and 'Latest stream ARN' as '-'. There is a 'Manage Stream' button. The 'Table details' section shows 'Table name' as 'GameScoresByUser', 'Primary partition key' as 'Username (String)', 'Primary sort key' as '-', and 'Point-in-time recovery' as 'DEFAULT'. The footer includes 'Feedback', 'English (US)', copyright information, 'Privacy Policy', 'Terms of Use', and a system clock showing 1:51 AM on 8/9/2019.

Create table Delete table

Filter by table name

Choose a table ... Actions

Name

- GameScoreRecords
- GameScoresByUser

GameScoresByUser Close

Overview Items Metrics Alarms Capacity Indexes Global Tables Backups Triggers Access control Tags

Table is being created

Recent alerts

No CloudWatch alarms have been triggered for this table.

Stream details

Stream enabled No

View type -

Latest stream ARN -

Manage Stream

Table details

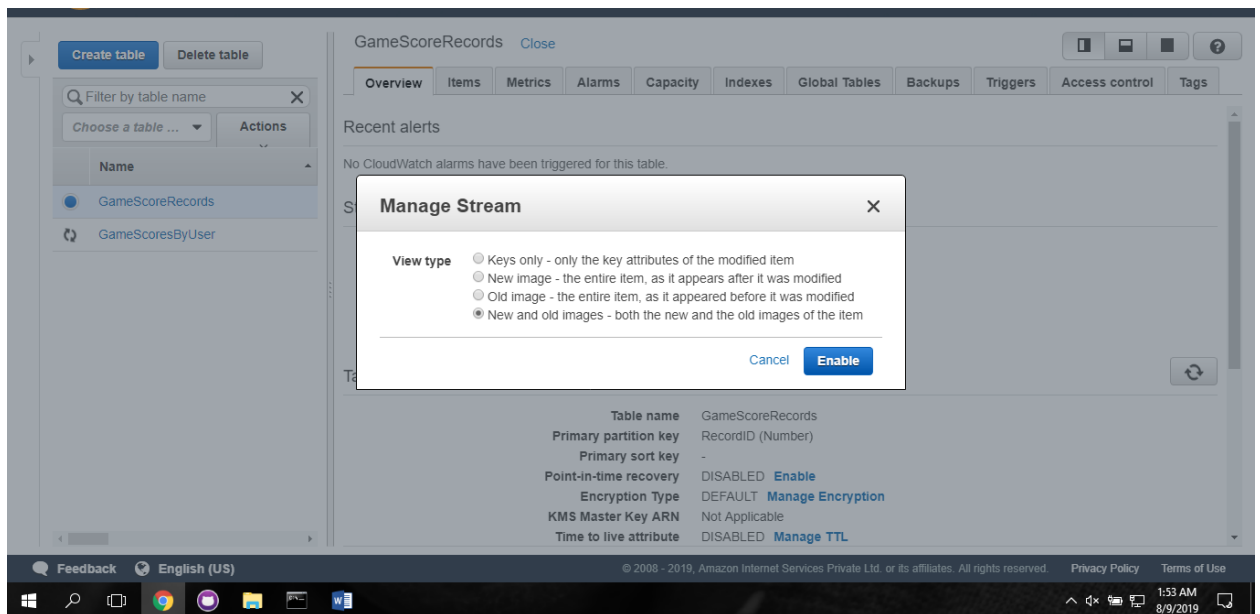
Table name GameScoresByUser

Primary partition key Username (String)

Primary sort key -

Point-in-time recovery DEFAULT

Feedback English (US) © 2008 - 2019, Amazon Internet Services Private Ltd. or its affiliates. All rights reserved. Privacy Policy Terms of Use 1:51 AM 8/9/2019



After the tables are ready, select the first table and click on Manage Stream from the Overview tab of the table. Select the appropriate option and click on Enable.

# Create a Lambda Function

Edit the previously created IAM Role, and attach AmazonDynamoDBFullAccess Policy.

Create a Lambda Function.

**Basic information**

**Function name**  
Enter a name that describes the purpose of your function.

Use only letters, numbers, hyphens, or underscores with no spaces.

**Runtime** [Info](#)  
Choose the language to use to write your function.

Node.js 8.10

**Permissions** [Info](#)  
Lambda will create an execution role with permission to upload logs to Amazon CloudWatch Logs. You can configure and modify permissions further when you add triggers.

▼ Choose or create an execution role

**Execution role**  
Choose a role that defines the permissions of your function. To create a custom role, go to the [IAM console](#).

Use an existing role

**Existing role**  
Choose an existing role that you've created to be used with this Lambda function. The role must have permission to upload logs to Amazon CloudWatch Logs.

LambdaRole

[View the LambdaRole role](#) on the IAM console.

Cancel

Create function


Scroll down to the Code editor. In the index.js part, paste the following code:

<https://pastebin.com/TV8C9r6j>

Click on Add Trigger button.

### Add trigger

#### Trigger configuration

 **DynamoDB**  
aws database nosql

**DynamoDB table**  
Select a DynamoDB table to listen for updates on.  

GameScoreRecords

**Batch size**  
The largest number of records that will be read from your table's update stream at once.  

100

**Starting position**  
The position in the stream to start reading from. For more information, see [ShardIteratorType](#) in the Amazon DynamoDB Streams API Reference.  

Latest

In order to read from the DynamoDB trigger, your execution role must have proper permissions.

☒ **Enable trigger**  
Enable the trigger now, or create it in a disabled state for testing (recommended).

Cancel


Add

Click on Test button.

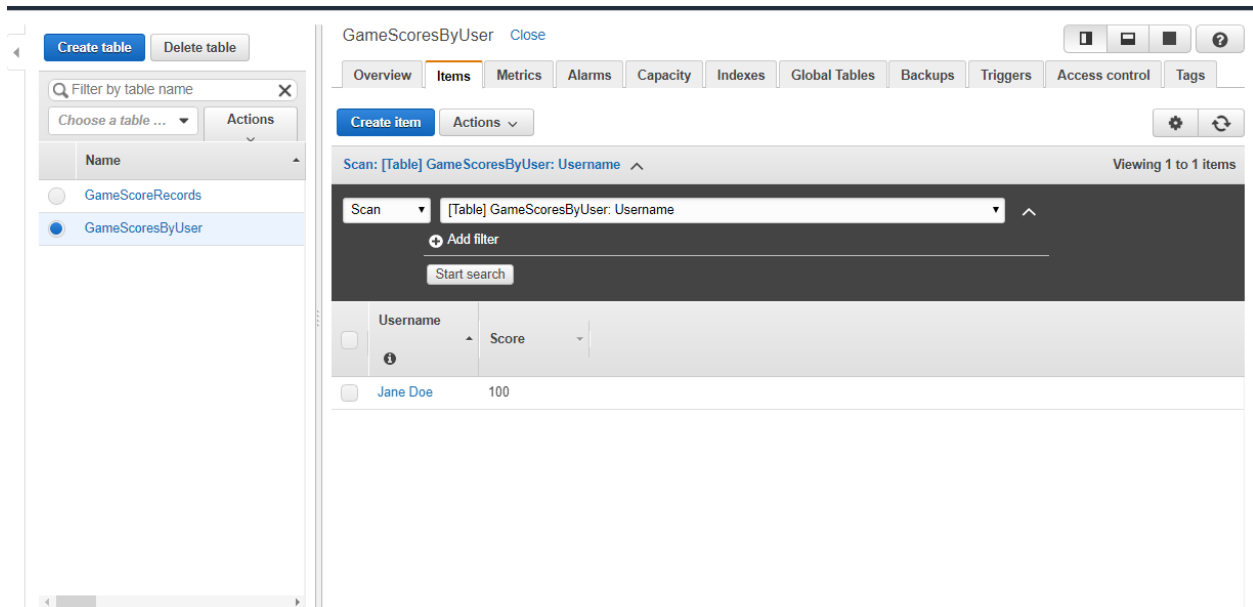
In the code editor that appears, paste the following code:

<https://pastebin.com/cw0jTEgi>

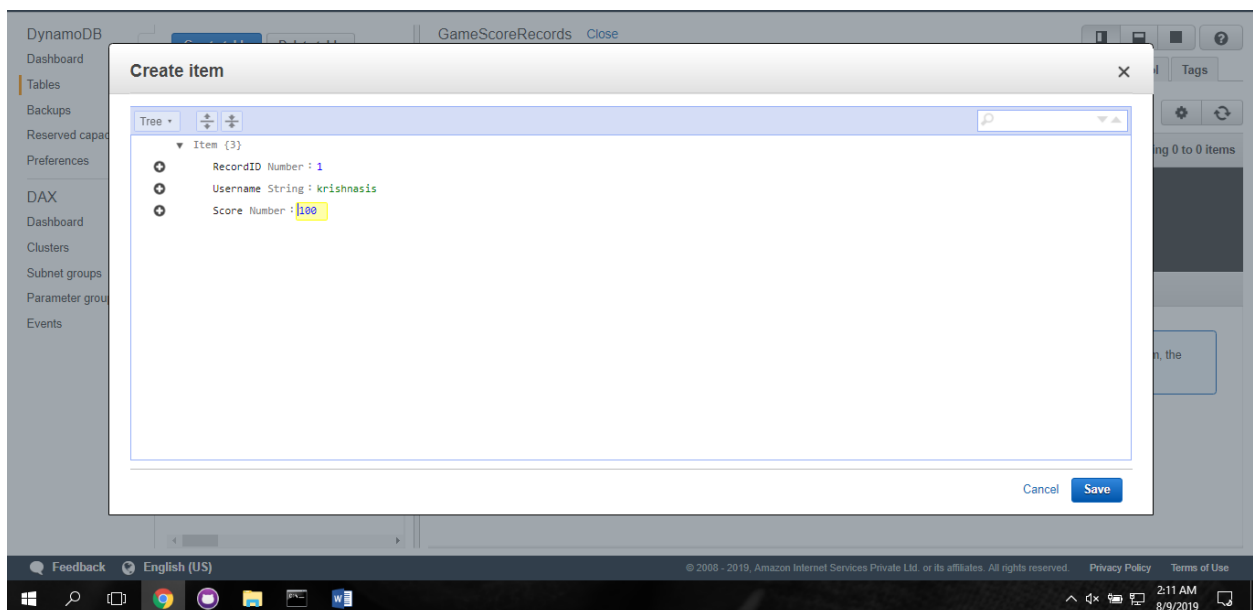
Click on the Test button again.



Navigate to the DynamoDB tables and click on the table that we entered the test data into. You will see the data is reflecting in the table.



Create an item in the other table, to test if the Lambda function is working or not.



Navigate to the other table, the record will reflect there as well. This shows that the Lambda function works properly.

Create tableDelete table

Filter by table name X

Choose a table ...Actions

Name

GameScoreRecords

GameScoresByUser

GameScoresByUserClose

OverviewItemsMetricsAlarmsCapacityIndexesGlobal TablesBackupsTriggersAccess controlTags

Create itemActions

Scan: [Table] GameScoresByUser: Username ^Viewing 1 to 2 items

Scan [Table] GameScoresByUser: Username ^

+ Add filter

Start search

UsernameScore

Jane Doe100

krishnasis100

# Thank You